

Regression/Classification Models

2023-12-07

Contents

```
# import train and test data
train_data = read.csv("data/condensed_train_data.csv")
test_data = read.csv("data/condensed_test_data.csv")
```

```
# X is just an index, remove it
train_data <- subset(train_data, select = -X)
test_data <- subset(test_data, select = -X)
```

```
head(train_data)
```

```
##           id total_events total_nonproduction_events total_remove_cut_events
## 1 001519c8         2557             120                417
## 2 0042269b         4136             175                439
## 3 0059420b         1556              99                151
## 4 0075873a         2531              72                517
## 5 0093f095         1765              34                148
## 6 009e23ab         2353             155                222
## total_paste_events total_replace_events avg_action_time max_action_time
## 1              0              7      116.24677      2259
## 2              0              7      101.83777      3005
## 3              1              1      121.84833       806
## 4              0              0      123.94390       701
## 5              0              0      109.71785       501
## 6              0              1      90.75563       803
## min_action_time sd_action_time final_word_count avg_text_change_length
## 1              0      91.79737          256      1.370747
## 2              0      82.38377          404      1.422872
## 3              0     113.76823          206      1.462725
## 4              0      62.08201          252      1.199131
## 5              0      37.01833          242      1.134844
## 6              0      41.93495          308      1.463238
## total_text_removed max_cursor_movement avg_cursor_movement sd_cursor_movement
## 1              524          591      0.4092332      43.37815
## 2              970         1826      0.1187424      72.08293
## 3              168          100      0.5125402      10.01403
## 4              517          468      0.5541502      24.40666
## 5              148          223      0.8134921      11.11349
## 6              228         1613      0.7444728     102.45500
## score
## 1    3.5
```

```
## 2 6.0
## 3 2.0
## 4 4.0
## 5 4.5
## 6 4.0
```

```
head(test_data)
```

```
##      id total_events total_nonproduction_events total_remove_cut_events
## 1 0022f953      2454              254              260
## 2 0081af50      2211              76              338
## 3 00e1f05a      7826             228             1446
## 4 0190ff4c      1922              46              118
## 5 01c359fc      2934             155              251
## 6 01d602a7      3573              73              358
##  total_paste_events total_replace_events avg_action_time max_action_time
## 1              1              1      112.22127      1758
## 2              0              3      81.40434      1102
## 3              0              7      93.34321     11017
## 4              0              0      98.37773       219
## 5              0              0     126.43626       582
## 6              0              0     145.40862       501
##  min_action_time sd_action_time final_word_count avg_text_change_length
## 1              0      55.43119             323      1.729014
## 2              0      40.65305             275      1.251922
## 3              0     198.89669             739      1.347048
## 4              0      28.53198             299      1.167534
## 5              0      46.29272             430      1.369802
## 6              0      50.09804             487      1.143017
##  total_text_removed max_cursor_movement avg_cursor_movement sd_cursor_movement
## 1              271             1336      0.6192417      85.350330
## 2              366             477      0.6538462      28.833439
## 3             2573            2668      0.5252396      73.546520
## 4              118             513      0.4232171      36.891782
## 5              251            1970      0.6137061      87.148047
## 6              358              33      0.7793953      1.160606
##  score
## 1 3.5
## 2 2.0
## 3 4.5
## 4 4.0
## 5 3.5
## 6 4.5
```

```
dim(train_data)
```

```
## [1] 1976 17
```

```
dim(test_data)
```

```
## [1] 495 17
```

```
train_x = train_data[, 2:16]
train_y = train_data[, 17]

test_x = test_data[, 2:16]
test_y = test_data[, 17]
```

We will start with elastic net regression, to do both feature selection and address collinearity issues that may exist within the data.

```
set.seed(432)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
# Grid of alpha values to try
# alpha_values <- c(0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1)
alpha_values <- seq(0, 1, by = 0.01)

# Create an empty matrix to store results
results <- matrix(NA, nrow = length(alpha_values), ncol = 2, dimnames = list(NULL, c("Alpha", "RMSE")))

# Perform grid search
for (i in seq_along(alpha_values)) {
  alpha <- alpha_values[i]
  lasso.fit <- cv.glmnet(data.matrix(train_x), train_y, alpha = alpha)

  # Find the index of the lambda that minimizes CV error
  min_lambda_index <- which.min(lasso.fit$cvm)

  # Optimal lambda and corresponding RMSE
  optimal_lambda <- lasso.fit$lambda[min_lambda_index]
  predictions <- predict(lasso.fit, newx = data.matrix(test_x), s = optimal_lambda)
  rmse <- sqrt(mean((predictions - test_y)^2))

  # Store results
  results[i, ] <- c(alpha, rmse)
}

# Find the row with the minimum RMSE
min_rmse_row <- which.min(results[, "RMSE"])

# Optimal alpha and lambda
optimal_alpha <- results[min_rmse_row, "Alpha"]
optimal_lambda <- lasso.fit$lambda[which.min(lasso.fit$cvm)]

# Print optimal values
cat("Optimal Alpha:", optimal_alpha, "\n")
```

```
## Optimal Alpha: 0.98
```

```
cat("Optimal Lambda:", optimal_lambda, "\n")
```

```
## Optimal Lambda: 0.002004471
```

```
cat("Optimal RMSE:", results[min_rmse_row, "RMSE"], "\n")
```

```
## Optimal RMSE: 0.7546587
```

As we can see, at the values $\alpha = 0.98$, $\lambda = 0.002004471$ return the optimal RMSE for the elastic net models, with an RMSE = 0.7546587

```
# Assuming elastic_net_fit is your fitted Elastic Net model  
elastic_net_fit <- cv.glmnet(data.matrix(train_data[, 2:16]), train_data[, 17], alpha = 0.98)
```

```
# Get coefficients for the optimal lambda  
optimal_lambda <- 0.002004471
```

```
coefficients <- coef(elastic_net_fit, s = optimal_lambda)
```

```
# Print or inspect the coefficients  
print(coefficients)
```

```
## 16 x 1 sparse Matrix of class "dgCMatrix"  
##               s1  
## (Intercept)    2.285071e+00  
## total_events   3.845065e-04  
## total_nonproduction_events -1.903365e-04  
## total_remove_cut_events -4.182445e-04  
## total_paste_events -1.161791e-03  
## total_replace_events  1.840437e-02  
## avg_action_time -1.252771e-04  
## max_action_time  1.594265e-06  
## min_action_time -4.849160e-02  
## sd_action_time    .  
## final_word_count  1.049640e-03  
## avg_text_change_length -9.385741e-02  
## total_text_removed -1.689733e-04  
## max_cursor_movement .  
## avg_cursor_movement -7.187129e-02  
## sd_cursor_movement  5.070199e-03
```

```
order(abs(coefficients), decreasing = TRUE)
```

```
## <sparse>[ <logic> ]: .M.sub.i.logical() maybe inefficient
```

```
## [1]  1 12 15  9  6 16  5 11  4  2  3 13  7  8 10 14
```

```
rownames(coefficients)[order(abs(coefficients), decreasing = TRUE)]
```

```
## <sparse>[ <logic> ]: .M.sub.i.logical() maybe inefficient
```

```
## [1] "(Intercept)"          "avg_text_change_length"
## [3] "avg_cursor_movement"   "min_action_time"
## [5] "total_replace_events"  "sd_cursor_movement"
## [7] "total_paste_events"    "final_word_count"
## [9] "total_remove_cut_events" "total_events"
## [11] "total_nonproduction_events" "total_text_removed"
## [13] "avg_action_time"       "max_action_time"
## [15] "sd_action_time"        "max_cursor_movement"
```

As we can see from the coefficient output, the elastic net model removed the variables “sd_action_time” and “max_cursor_movement” indicating that they may not be important for predicting the score of the row. Noting the largest coefficients by magnitude, we note that the variables “avg_text_change_length”, “avg_cursor_movement”, “min_action_time” seem important for predicting the score.

We will now try k-nearest neighbors

```
set.seed(432)

# "1-nearest neighbor" regression using kkn package
library(kknn)
```

```
## Warning: package 'kknn' was built under R version 4.3.2
```

```
knn.fit = kknn(y ~ ., train = data.frame(x = train_data[, 2:16], y = train_data[, 17]),
             test = data.frame(x = test_data[, 2:16]),
             k = 1, kernel = "rectangular")
test.pred = knn.fit$fitted.values
```

```
# Calculate the root mean squared error (RMSE)
rmse <- sqrt(mean((test.pred - test_y)^2))

# Print or use the RMSE
cat("Root Mean Squared Error (RMSE):", rmse, "\n")
```

```
## Root Mean Squared Error (RMSE): 1.028876
```

```
library(kknn)

set.seed(432)

# Specify a range of k values to test
k_values <- seq(1, 200, by = 1) # Adjust the range as needed

# Initialize variables to store results
rmse_values <- numeric(length(k_values))

# Loop over different k values
for (i in seq_along(k_values)) {
  # Fit KNN model
  knn_fit <- kknn(y ~ ., train = data.frame(x = train_x, y = train_y),
```

```

        test = data.frame(x = test_x),
        k = k_values[i])

# Make predictions
test_pred <- knn_fit$fitted.values

# Calculate RMSE
residuals <- test_pred - test_y
mse <- mean(residuals^2)
rmse_values[i] <- sqrt(mse)
}

k_values

```

```

##      [1]      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16     17     18
##     [19]     19     20     21     22     23     24     25     26     27     28     29     30     31     32     33     34     35     36
##     [37]     37     38     39     40     41     42     43     44     45     46     47     48     49     50     51     52     53     54
##     [55]     55     56     57     58     59     60     61     62     63     64     65     66     67     68     69     70     71     72
##     [73]     73     74     75     76     77     78     79     80     81     82     83     84     85     86     87     88     89     90
##     [91]     91     92     93     94     95     96     97     98     99    100    101    102    103    104    105    106    107    108
##    [109]    109    110    111    112    113    114    115    116    117    118    119    120    121    122    123    124    125    126
##    [127]    127    128    129    130    131    132    133    134    135    136    137    138    139    140    141    142    143    144
##    [145]    145    146    147    148    149    150    151    152    153    154    155    156    157    158    159    160    161    162
##    [163]    163    164    165    166    167    168    169    170    171    172    173    174    175    176    177    178    179    180
##    [181]    181    182    183    184    185    186    187    188    189    190    191    192    193    194    195    196    197    198
##    [199]    199    200

```

```
rmse_values
```


```

##      [1] 1.0288760 0.9549758 0.9025537 0.8730921 0.8531476 0.8381311 0.8269766
##      [8] 0.8186804 0.8123432 0.8067794 0.8021411 0.7983028 0.7948448 0.7917984
##     [15] 0.7890232 0.7864230 0.7839734 0.7816491 0.7795199 0.7775846 0.7758163
##     [22] 0.7742726 0.7729047 0.7716265 0.7704527 0.7693762 0.7683433 0.7673268
##     [29] 0.7663343 0.7654100 0.7645595 0.7637800 0.7630919 0.7624803 0.7619254
##     [36] 0.7614086 0.7609268 0.7604792 0.7600553 0.7596652 0.7593228 0.7589931
##     [43] 0.7586347 0.7582742 0.7579403 0.7576231 0.7573123 0.7570240 0.7567658
##     [50] 0.7565206 0.7562866 0.7560785 0.7558970 0.7557252 0.7555520 0.7553841
##     [57] 0.7552202 0.7550559 0.7548916 0.7547329 0.7545797 0.7544280 0.7542872
##     [64] 0.7541491 0.7540156 0.7538965 0.7537818 0.7536648 0.7535507 0.7534454
##     [71] 0.7533521 0.7532740 0.7532145 0.7531655 0.7531129 0.7530626 0.7530167
##     [78] 0.7529718 0.7529318 0.7529054 0.7528932 0.7528783 0.7528619 0.7528502
##     [85] 0.7528404 0.7528356 0.7528342 0.7528376 0.7528442 0.7528505 0.7528566
##     [92] 0.7528587 0.7528585 0.7528593 0.7528622 0.7528681 0.7528788 0.7528984
##     [99] 0.7529231 0.7529468 0.7529727 0.7530016 0.7530310 0.7530605 0.7530916
##    [106] 0.7531267 0.7531635 0.7532006 0.7532397 0.7532797 0.7533159 0.7533514
##    [113] 0.7533930 0.7534406 0.7534918 0.7535461 0.7536054 0.7536688 0.7537352
##    [120] 0.7538043 0.7538755 0.7539474 0.7540183 0.7540864 0.7541522 0.7542194
##    [127] 0.7542882 0.7543571 0.7544262 0.7544970 0.7545676 0.7546368 0.7547074
##    [134] 0.7547794 0.7548516 0.7549235 0.7549943 0.7550651 0.7551364 0.7552086
##    [141] 0.7552832 0.7553595 0.7554383 0.7555192 0.7555994 0.7556800 0.7557604
##    [148] 0.7558400 0.7559205 0.7560000 0.7560793 0.7561574 0.7562348 0.7563153
##    [155] 0.7563975 0.7564785 0.7565577 0.7566371 0.7567169 0.7567968 0.7568764
##    [162] 0.7569567 0.7570384 0.7571213 0.7572047 0.7572879 0.7573706 0.7574532

```

```
## [169] 0.7575348 0.7576155 0.7576966 0.7577779 0.7578601 0.7579432 0.7580263
## [176] 0.7581099 0.7581950 0.7582806 0.7583654 0.7584506 0.7585359 0.7586192
## [183] 0.7587022 0.7587858 0.7588686 0.7589518 0.7590352 0.7591166 0.7591958
## [190] 0.7592743 0.7593522 0.7594291 0.7595062 0.7595840 0.7596619 0.7597405
## [197] 0.7598195 0.7598987 0.7599778 0.7600565
```

```
# want to plot rmse vs k_values
plot(k_values, rmse_values)
```



```
# Find the k that gives the lowest RMSE
optimal_k <- k_values[which.min(rmse_values)]

# Print results
cat("Optimal K:", optimal_k, "\n")
```

```
## Optimal K: 87
```

```
cat("Minimum RMSE:", min(rmse_values), "\n")
```

```
## Minimum RMSE: 0.7528342
```

We see that the optimal $k = 87$ returns the lowest RMSE for the k -nearest neighbors regression with $\text{RMSE} = 0.7528342$.

We will now try random forest.

```
set.seed(432)
```

```
# fit random forests with a selected tuning  
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 4.3.2
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
rf.fit = randomForest(train_x, train_y)
```

```
summary(rf.fit)
```

```
##           Length Class  Mode  
## call           3  -none- call  
## type           1  -none- character  
## predicted     1976  -none- numeric  
## mse           500  -none- numeric  
## rsq           500  -none- numeric  
## oob.times     1976  -none- numeric  
## importance     15  -none- numeric  
## importanceSD    0  -none- NULL  
## localImportance 0  -none- NULL  
## proximity      0  -none- NULL  
## ntree          1  -none- numeric  
## mtry           1  -none- numeric  
## forest         11  -none- list  
## coefs          0  -none- NULL  
## y             1976  -none- numeric  
## test          0  -none- NULL  
## inbag          0  -none- NULL
```

```
# ntree = 500  
# sampsize = ?  
# mtry = 5  
# nodesize ?
```

```
# Assuming 'test_x' contains your test features and 'test_y' contains your test target variable
```

```
# Make predictions on the test data
```

```
rf_predictions <- predict(rf.fit, newdata = test_x)
```

```
# Calculate the root mean squared error (RMSE)
```

```
rmse <- sqrt(mean((rf_predictions - test_y)^2))
```

```
# Print or use the RMSE
```

```
cat("Root Mean Squared Error (RMSE):", rmse, "\n")
```

```
## Root Mean Squared Error (RMSE): 0.6985063
```


Random forest has the best RMSE so far, with an $RMSE = 0.6985063$.

```
library(randomForest)

# Set seed for reproducibility
set.seed(432)

# Define candidate values for parameters
ntree_values <- c(400, 500, 600, 700, 800)
sampsize_values <- c(90, 100, 110, 120, 130)
mtry_values <- c(12, 13, 14, 15)
nodesize_values <- c(1, 2, 3, 4, 5)

# Initialize variables to store optimal values
best_rmse <- Inf
optimal_ntree <- NULL
optimal_sampsize <- NULL
optimal_mtry <- NULL
optimal_nodesize <- NULL
best_model <- NULL

# Iterate over parameter combinations
for (ntree in ntree_values) {
  for (sampsize in sampsize_values) {
    for (mtry in mtry_values) {
      for (nodesize in nodesize_values) {

        # Train the model
        rf_model <- randomForest(train_y ~ ., data = train_x,
                                ntree = ntree, mtry = mtry,
                                sampsize = sampsize, nodesize = nodesize)

        # Make predictions on the test data
        rf_predictions <- predict(rf_model, newdata = test_x)

        # Calculate RMSE
        rmse <- sqrt(mean((rf_predictions - test_y)^2))

        # Check if current combination improves RMSE
        if (rmse < best_rmse) {
          best_rmse <- rmse
          optimal_ntree <- ntree
          optimal_sampsize <- sampsize
          optimal_mtry <- mtry
          optimal_nodesize <- nodesize
          best_model <- rf_model
        }
      }
    }
  }
}

# Print or use the optimal values
cat("Optimal ntree:", optimal_ntree, "\n")
```

```
## Optimal ntree: 700
```

```
cat("Optimal sampsize:", optimal_sampsize, "\n")
```

```
## Optimal sampsize: 110
```

```
cat("Optimal mtry:", optimal_mtry, "\n")
```

```
## Optimal mtry: 13
```

```
cat("Optimal nodesize:", optimal_nodesize, "\n")
```

```
## Optimal nodesize: 1
```

```
cat("Optimal Root Mean Squared Error (RMSE):", best_rmse, "\n")
```

```
## Optimal Root Mean Squared Error (RMSE): 0.6829625
```

As we can see from our gridsearch, the best model has ntree=700, sampsize=110, mtry=13, nodesize=1, with RMSE = 0.6829625.

```
best_model$importance
```

```
##                               IncNodePurity
## total_events                  17.4698577
## total_nonproduction_events    3.7838847
## total_remove_cut_events       3.9744701
## total_paste_events            0.4763130
## total_replace_events          1.8217903
## avg_action_time               4.4580144
## max_action_time               4.0362660
## min_action_time               0.3842305
## sd_action_time                3.6217054
## final_word_count              51.0266429
## avg_text_change_length        4.1440921
## total_text_removed            3.9363097
## max_cursor_movement           5.3090077
## avg_cursor_movement           4.5206313
## sd_cursor_movement            4.7586477
```

```
order(best_model$importance, decreasing = TRUE)
```

```
## [1] 10  1 13 15 14  6 11  7  3 12  2  9  5  4  8
```

```
rownames(best_model$importance)[order(best_model$importance, decreasing = TRUE)]
```

```
## [1] "final_word_count"          "total_events"
## [3] "max_cursor_movement"      "sd_cursor_movement"
## [5] "avg_cursor_movement"      "avg_action_time"
## [7] "avg_text_change_length"   "max_action_time"
## [9] "total_remove_cut_events"   "total_text_removed"
## [11] "total_nonproduction_events" "sd_action_time"
## [13] "total_replace_events"      "total_paste_events"
## [15] "min_action_time"
```

From the variable importance we see that the variables “final_word_count”, “total_events”, and “max_cursor_movement” are important in predicting the outcome in the random forest model.

We will now try Support Vector Machine classification.

```
set.seed(432)

library(e1071)

# Assuming 'train_data' contains your training data with an ordinal target variable
# Replace 'target' with the actual column name of your target variable

# Fit an SVM regression model
svm_model <- svm(as.factor(train_y) ~ ., data = train_x, kernel = "linear", type = "C-classification", cost = 1)

# Make predictions on new data
predictions <- predict(svm_model, newdata = test_x)

# Evaluate the model's performance as needed
# Calculate the confusion matrix
conf_matrix <- table(Actual = test_y, Predicted = predictions)

# Calculate the RMSE
rmse <- sqrt(mean((as.numeric(predictions) - test_y)^2))

# Print the confusion matrix
print("Confusion Matrix:")
```

```
## [1] "Confusion Matrix:"
```

```
print(conf_matrix)
```

```
##          Predicted
## Actual 0.5  1  1.5  2  2.5  3  3.5  4  4.5  5  5.5  6
##   0.5    0  0   0  0   0   1   0  0   0  0   0  0
##   1      0  0   0  0   0   6   0  0   0  0   0  0
##   1.5    0  0   0  0   0   9   5  2   0  0   0  0
##   2      0  0   0  0   0  15   6  0   0  0   0  0
##   2.5    0  0   0  0   0  33  10  3   0  0   0  0
##   3      0  0   0  0   1  31  19 10   4  0   0  0
##   3.5    0  0   0  0   0  29  41 24   8  0   0  0
##   4      0  0   0  0   0  14  22 37  17  0   1  0
##   4.5    0  0   0  0   0   2  10 29  41  0   0  0
```

```
##      5      0 0 0 0 0 1 2 6 19 0 2 0
##     5.5    0 0 0 0 0 0 1 4 21 0 3 0
##      6      0 0 0 0 0 0 0 2 3 0 1 0
```

```
accuracy = sum(predictions == test_y) / length(test_y)
cat("Accuracy:", accuracy, "\n")
```

```
## Accuracy: 0.3090909
```

```
# Print the RMSE
cat("Root Mean Squared Error (RMSE):", rmse, "\n")
```

```
## Root Mean Squared Error (RMSE): 3.898977
```

```
set.seed(432)

library(e1071)

# Assuming 'train_data' and 'test_data' contain your training and testing data

# Define the values of cost to try
cost_values <- c(1, 2,3,4,5, 10, 50, 100)

# Initialize a vector to store RMSE values
rmse_values <- numeric(length(cost_values))

# Fit SVM models with different cost values
for (i in seq_along(cost_values)) {
  cost <- cost_values[i]

  # Fit the SVM model
  svm_model <- svm(as.factor(train_y) ~ ., data = train_x, kernel = "linear", cost = cost)

  # Make predictions on the test data
  predictions <- predict(svm_model, newdata = test_x)

  # Calculate RMSE
  rmse <- sqrt(mean((as.numeric(predictions) - test_y)^2))

  # Store the RMSE value
  rmse_values[i] <- rmse

  # Print or save other information as needed
  cat("Cost:", cost, " - RMSE:", rmse, "\n")

  # Assuming 'test_data' contains your testing data with 'target_class' as the true class

  # Calculate accuracy
  accuracy <- sum(predictions == test_y) / length(test_y)

  # Print the accuracy
  cat("Accuracy:", accuracy, "\n")
}
```

```
## Cost: 1 - RMSE: 3.898977
## Accuracy: 0.3090909
## Cost: 2 - RMSE: 3.925313
## Accuracy: 0.3070707
## Cost: 3 - RMSE: 3.928914
## Accuracy: 0.3070707
## Cost: 4 - RMSE: 3.925313
## Accuracy: 0.3070707
## Cost: 5 - RMSE: 3.929942
## Accuracy: 0.3090909
## Cost: 10 - RMSE: 3.948405
## Accuracy: 0.3131313
## Cost: 50 - RMSE: 3.935079
## Accuracy: 0.3111111
## Cost: 100 - RMSE: 3.93097
## Accuracy: 0.3111111
```

```
# Find the optimal cost value
optimal_cost <- cost_values[which.min(rmse_values)]
cat("Optimal Cost:", optimal_cost, "\n")
```

```
## Optimal Cost: 1
```

As we can see from the output, the optimal cost = 1, with an RMSE = 3.898977. What is interesting to note is that even though cost=10 has a higher RMSE, it also has better classification accuracy.

SVM Regression

```
set.seed(432)

library(e1071)

# Assuming 'train_data' and 'test_data' contain your training and testing data

# Define the values of cost to try
cost_values <- c(1, 30, 50, 70, 75, 80, 100)

# Initialize a vector to store RMSE values
rmse_values <- numeric(length(cost_values))

# Fit SVM models with different cost values
for (i in seq_along(cost_values)) {
  cost <- cost_values[i]

  # Fit the SVM model
  svm_model <- svm(train_y ~ ., data = train_x, kernel = "linear", cost = cost)

  # Make predictions on the test data
  predictions <- predict(svm_model, newdata = test_x)

  # Calculate RMSE
  rmse <- sqrt(mean((predictions - test_y)^2))
}
```

```

# Store the RMSE value
rmse_values[i] <- rmse

# Print or save other information as needed
cat("Cost:", cost, " - RMSE:", rmse, "\n")

# Assuming 'test_data' contains your testing data with 'target_class' as the true class
}

```

```

## Cost: 1 - RMSE: 0.7603172
## Cost: 30 - RMSE: 0.7601203
## Cost: 50 - RMSE: 0.760139
## Cost: 70 - RMSE: 0.7601615
## Cost: 75 - RMSE: 0.7580726
## Cost: 80 - RMSE: 0.7597441
## Cost: 100 - RMSE: 0.7607451

```

```

# Find the optimal cost value
optimal_cost <- cost_values[which.min(rmse_values)]
cat("Optimal Cost:", optimal_cost, "\n")

```

```
## Optimal Cost: 75
```

```

optimal_rmse <- min(rmse_values)
cat("Optimal RMSE:", optimal_rmse, "\n")

```

```
## Optimal RMSE: 0.7580726
```

As we can see, SVM regression returns a much better RMSE = 0.7580726 with cost = 75 than SVM for classification.

```

models = c("elastic net", "k nearest neighbors", "random forests", "svm classification", "svm regression")
rmse_values = c(0.7546587, 0.7528342, 0.6829625, 3.898977, 0.7580726)

```

```

rmse_table = data.frame(Model = models, RMSE = rmse_values)
rmse_table

```

```

##           Model      RMSE
## 1 elastic net 0.7546587
## 2 k nearest neighbors 0.7528342
## 3 random forests 0.6829625
## 4 svm classification 3.8989770
## 5 svm regression 0.7580726

```

```

# sorted order
rmse_table[order(rmse_table$RMSE), ]

```

```

##           Model      RMSE
## 3 random forests 0.6829625
## 2 k nearest neighbors 0.7528342
## 1 elastic net 0.7546587
## 5 svm regression 0.7580726
## 4 svm classification 3.8989770

```